

OpenMP SIMD directives

```
void v_add (float *c, float *a, float *b)
{
#pragma omp simd
    for (int i = 0; i <= MAX; i++)
        c[i] = a[i] + b[i];
}
```

Vectorization solutions

1. auto-vectorization (use a compiler switch and hope it vectorizes)
2. give your compiler hints and hope it vectorizes
3. code explicitly

openMP 4.0 #pragma omp simd

CilkTM Plus array notations

SIMD instruction intrinsics

Kernels: OpenMP 4.0 #pragma omp declare

Optimization Reports

-opt-report

-opt-report-file=[<name>|stdout|stderr]

“Vectorization Advisor” – Advisor XE

Application Tuning – Amplifier XE

Questions

1. What is the difference between parallelism and vectorization? Can we think vectorization as a kind of parallelism?
2. What would the compiler do if my loop or the length of a vector is larger than the cache line? Will it break the vector into small parts automatically?

Answers

1. No. Because $a[i]$ and $b[i]$ are not necessarily in different places of memory than $c[i]$, every time an element is written into $c[i]$ it may have changed an element of $a[i]$ or $b[i]$. The original program was written to do the addition once at a time, so the compiler needs to take that completely literally.
2. Yes. The `declare SIMD` construct can be applied to a function to enable the creation of one or more versions that can process multiple arguments using SIMD instructions from a single invocation from a SIMD loop.